

Module 2 Lab Exercise: Tools Used in Machine Learning

Learning Objectives

By the end of this lab, you will be able to:

- Set up and navigate Jupyter Notebook, Google Colab, and VS Code environments
- Install and import essential Python libraries for machine learning
- Create and format professional documentation using Markdown
- Initialize a GitHub repository for your ML projects
- Understand the basic workflow of data science tools

Prerequisites

- Basic understanding of what machine learning is (Module 1)
 - Access to internet for downloading tools and datasets
 - A Google account (for Colab) or local Python installation
-

Part 1 : Environment Setup and Tool Overview

What are the main tools we'll use in this course?

Jupyter Notebook/Google Colab: Interactive computing environments where you can write code, see results immediately, and document your work with text and visualizations.

Python Libraries: Pre-written code packages that make machine learning tasks easier:

- Pandas: For working with data (like Excel, but more powerful)
- NumPy: For mathematical operations on arrays of numbers
- Matplotlib: For creating charts and graphs
- Scikit-learn: The main library for machine learning algorithms

GitHub: A platform to store, share, and collaborate on code projects

VS Code: A powerful text editor for writing and debugging code

Let's start by setting up our environment!

Environment Setup Instructions

Option 1 : Google Colab (Recommended for Beginners)

- 1 . Go to colab.research.google.com
- 2 . Sign in with your Google account
- 3 . Click "New Notebook"
- 4 . You're ready to go! Libraries are pre-installed.

Option 2 : Local Jupyter Notebook

- 1 . Install Python from python.org
- 2 . Open terminal/command prompt
- 3 . Run: `pip install jupyter pandas numpy matplotlib scikit-learn`
- 4 . Run: `jupyter notebook`
- 5 . Create a new notebook

Option 3 : VS Code

- 1 . Download VS Code from code.visualstudio.com
- 2 . Install Python extension
- 3 . Install Jupyter extension
- 4 . Create a new .ipynb file

For this lab, we recommend starting with Google Colab as it requires no installation.

```
In [1]: # Install required libraries (uncomment if needed)
# !pip install pandas numpy matplotlib scikit-learn

# Import libraries with standard aliases
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
import warnings
warnings.filterwarnings('ignore') # Hide warning messages for cleaner output

print("✅ All libraries imported successfully!")
print(f"Pandas version: {pd.__version__}")
print(f"NumPy version: {np.__version__}")
```

✅ All libraries imported successfully!
Pandas version: 2.2.3
NumPy version: 2.3.5

Part 2 : Loading and Exploring Your First Dataset

We'll use the famous Iris dataset - a classic dataset for beginners. It contains measurements of iris flowers three different species.

```
In [2]: # Load the Iris dataset from scikit-learn.
from sklearn.datasets import load_iris

iris = load_iris()

# Check the dimensions and labels so we understand what was loaded.
print("Dataset loaded successfully!")
print(f"Dataset shape: {iris.data.shape}")
print(f"Features: {iris.feature_names}")
print(f"Target classes: {iris.target_names}")
```

Dataset loaded successfully!

Dataset shape: (150, 4)

Features: ['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']

Target classes: ['setosa' 'versicolor' 'virginica']

```
In [3]: # Put the numerical features into a pandas DataFrame so they are easier to work with.
df = pd.DataFrame(iris.data, columns=iris.feature_names)

# Convert the numeric target values into readable species names.
df['species'] = iris.target_names[iris.target]

print("First 5 rows of our dataset:")
print(df.head())

print("\nDataset info:")
print(df.info())
```

First 5 rows of our dataset:

	sepal length (cm)	sepal width (cm)	...	petal width (cm)	species
0	5.1	3.5	...	0.2	setosa
1	4.9	3.0	...	0.2	setosa
2	4.7	3.2	...	0.2	setosa
3	4.6	3.1	...	0.2	setosa
4	5.0	3.6	...	0.2	setosa

[5 rows x 5 columns]

Dataset info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 150 entries, 0 to 149

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	sepal length (cm)	150 non-null	float64
1	sepal width (cm)	150 non-null	float64
2	petal length (cm)	150 non-null	float64
3	petal width (cm)	150 non-null	float64
4	species	150 non-null	object

dtypes: float64(4), object(1)

memory usage: 6.0+ KB

None

Part 3 : Creating Your First Visualization

Data visualization is crucial in machine learning. Let's create a simple plot to understand our data.


```
In [4]: # Create a simple scatter plot
plt.figure(figsize=(10, 6))

# Plot sepal length vs sepal width, colored by species
species_colors = {'setosa': 'red', 'versicolor': 'blue', 'virginica': 'green'}

for species in df['species'].unique():
    species_data = df[df['species'] == species]
    plt.scatter(species_data['sepal length (cm)'],
                species_data['sepal width (cm)'],
                c=species_colors[species],
                label=species,
                alpha=0.7)

plt.xlabel('Sepal Length (cm)')
plt.ylabel('Sepal Width (cm)')
plt.title('Iris Dataset: Sepal Length vs Sepal Width')
plt.legend()
plt.grid(True, alpha=0.3)
plt.show()

print("🎉 Congratulations! You've created your first data visualization!")
```



🎉 Congratulations! You've created your first data visualization!

Part 4 : Practice with Basic Data Operations

Let's practice some basic data analysis operations that you'll use throughout the course.

```
In [5]: # Basic statistical analysis
print("Basic Statistics for Iris Dataset:")
print("=" * 40)

# Calculate mean values for each species
species_means = df.groupby('species').mean()
print("\nMean values by species:")
print(species_means)

# Count samples per species
species_counts = df['species'].value_counts()
print("\nSamples per species:")
print(species_counts)
```

Basic Statistics for Iris Dataset:

=====

Mean values by species:

	sepal length (cm)	...	petal width (cm)
species		...	
setosa	5.006	...	0.246
versicolor	5.936	...	1.326
virginica	6.588	...	2.026

[3 rows x 4 columns]

Samples per species:

```
species
setosa      50
versicolor  50
virginica   50
Name: count, dtype: int64
```

Part 5 : GitHub and Documentation Best Practices

Why GitHub for Machine Learning?

- Version Control: Track changes to your code and data
- Collaboration: Work with others on projects
- Portfolio: Showcase your work to potential employers
- Backup: Never lose your work

Basic GitHub Workflow:

- 1 . Create Repository: A folder for your project
- 2 . Clone/Download: Get the project on your computer
- 3 . Add Files: Put your notebooks and data
- 4 . Commit: Save a snapshot of your changes
- 5 . Push: Upload changes to GitHub

For This Course:

- Create a repository named "ITAI- 1 3 7 1 -ML-Labs"
- Upload each lab notebook as you complete it
- Include a README.md file describing your projects

Action Item: After this lab, create your GitHub account and repository.

Assessment: Tool Familiarity Check

Complete the following tasks to demonstrate your understanding of the tools:

```
In [6]: # Task 1: Use NumPy to calculate two descriptive statistics.
# The mean gives the average sepal length; the standard deviation describes the spread.

sepal_lengths = df['sepal length (cm)']

mean_sepal_length = np.mean(sepal_lengths)
std_sepal_length = np.std(sepal_lengths)

print(f"Mean sepal length: {mean_sepal_length:.2f} cm")
print(f"Standard deviation: {std_sepal_length:.2f} cm")

# Verification (don't modify)
assert isinstance(mean_sepal_length, (float, np.floating)), "Mean should be a float"
assert isinstance(std_sepal_length, (float, np.floating)), "Std should be a float"
print("✅ Task 1 completed successfully!")
```

Mean sepal length: 5.84 cm


Standard deviation: 0.83 cm

✅ Task 1 completed successfully!

```
In [7]: # Task 2: Create a simple bar chart showing species counts
species_counts = df['species'].value_counts()

plt.figure(figsize=(8, 5))
plt.bar(species_counts.index, species_counts.values, color=['red', 'blue'])
plt.title('Number of Samples per Species')
plt.xlabel('Species')
plt.ylabel('Count')
plt.show()

print(f"Species distribution: {dict(species_counts)}")
print("✅ Task 2 completed successfully!")
```



Species distribution: {'setosa': np.int64(50), 'versicolor': np.int64(50), 'virginica': np.int64(50)}

✅ Task 2 completed successfully!

Your Analysis and Reflection

My Observations About the Iris Dataset

Dataset Overview:

- Number of samples: 1 5 0
- Number of features: 4
- Number of classes: 3

Key Findings from the Visualization:

- 1 . The scatter plot shows that setosa flowers form a noticeably separate group from the other two species on sepal measurements.
- 2 . Versicolor and virginica overlap more than setosa, so sepal length and sepal width alone do not completely separate those two classes.
- 3 . The visualization makes it easier to recognize patterns that are harder to see from a table of numbers

Questions for Further Investigation:

- Would petal length and petal width separate the three species more clearly than sepal measurements?

- How would the results change if the data were standardized before analysis?

Reflection: Working with the Iris dataset helped me understand why data science tools are useful together than separately. Pandas made the dataset easier to organize and summarize, NumPy made numerical calculation straightforward, and Matplotlib helped me see patterns visually. I also learned that a visualization can reveal relationships and overlap that are not obvious when looking only at numerical output.

Lab Summary and Next Steps

What You've Accomplished:

- ✓ Set up your machine learning development environment
- ✓ Imported and used essential Python libraries
- ✓ Loaded and explored your first dataset
- ✓ Created your first data visualization
- ✓ Practiced professional documentation with Markdown
- ✓ Learned about GitHub for project management

Preparation for Module 3 :

In the next lab, you'll:

- Learn about different types of machine learning
- Build your first simple classifier
- Understand the complete ML workflow
- Work with more complex datasets

Action Items:

- 1 . Create your GitHub account and repository
- 2 . Upload this completed notebook to your repository
- 3 . Experiment with different visualizations using the Iris dataset
- 4 . Practice Markdown formatting in a new notebook

Resources for Continued Learning:

- [Pandas Documentation](#)
- [Matplotlib Gallery](#)
- [GitHub Guides](#)
- [Jupyter Notebook Tips](#)

Great job completing Module 2 ! You're now equipped with the essential tools for machine learning.